

Computer Technology and Parallelism

Computer Technology

- Performance improvements:
 - Improvements in semiconductor technology
 - Feature size, clock speed
 - Improvements in computer architectures
 - Enabled by HLL compilers, UNIX
 - Lead to RISC architectures
 - Improvement in cost-performance
 - Mobile devices, Personal computers and Workstations
- Together have enabled:
 - Lightweight computers
 - Enhanced the capability available to computer users
 - Software development: Productivity-based managed/interpreted programming languages (Java, C#, Python etc) instead of performance oriented languages (C and C++)
 - Change in the nature of applications: Speech, sound, images, video, text and many more unstructured data

2

Instruction Set Architectures

- **Complex Instruction Set Computer (CISC) processors:**
 - 2-operand instructions and 1-operand instructions
 - Any instruction can use memory operands
 - Many addressing modes
 - Complex instruction formats: Varying length instructions
 - Microprogrammed control unit
- **Reduced Instruction Set Computer (RISC) processors:**
 - 3-operand instructions, 2-operand instructions, and 1-operand instructions
 - **Load-Store Architecture (LSA) processors:**
 - Only memory transfer instructions (Load and Store) can use memory operands.
 - All other instructions can use register operands only.
 - A few addressing modes
 - Simple instruction formats: Fixed length instructions
 - Hardwired control unit
 - Suitable for pipelining


3

Single Processor Performance: Moore's Law

- The performance of computer hardware doubles every 1.5 years

Moore's Law – The number of transistors on integrated circuit chips (1971-2016) Our World in Data
 Moore's law describes the empirical regularity that the number of transistors on integrated circuits doubles approximately every two years. This advancement is important as other aspects of technological progress – such as processing speed or the price of electronic components – are strongly linked to Moore's law.

2016: 5.7×10^9



Date: 2016-03-01 10:10:10 (UTC+01:00) (Transistor count)
 The data visualization is available at ourworldindata.org, where you find more visualizations and research on this topic.

Licensed under CC-BY-SA by the author Max Roser.

4

Current Trends in Architecture

- Cannot continue to leverage Instruction-Level parallelism (ILP)
 - Single processor performance improvement ended in 2003
- New models for performance:
 - Data-level parallelism (DLP)
 - Thread-level parallelism (TLP)
 - Request-level parallelism (RLP)
- These require explicit restructuring of the application

5


Classes of Computers

- Personal Mobile Device (PMD)
 - e.g. smart phones, tablet computers
 - Emphasis on energy efficiency and real-time
- Desktop Computing
 - Emphasis on price-performance
- Servers
 - Emphasis on availability, scalability, throughput
- Clusters / Warehouse Scale Computers
 - Used for “Software as a Service (SaaS)”
 - Emphasis on availability and price-performance
 - Sub-class: Supercomputers, emphasis: floating-point performance and fast internal networks
- Internet of Things/Embedded Computers
 - Emphasis: price

6

Power

- Intel 80386 consumed ~ 2 W
- 3.3 GHz Intel Core i7 consumes 130 W
- Heat must be dissipated from 1.5×1.5 cm chip
- This is the limit of what can be cooled by air



7

Power and Energy

- Problem: Get power in, get power out
- Thermal Design Power (TDP)
 - Characterizes sustained power consumption
 - Used as target for power supply and cooling system
 - Lower than peak power (1.5X higher), higher than average power consumption
- Clock rate can be reduced dynamically to limit power consumption
- Reducing clock rate reduces power, not energy

8

Classes of Parallelism and Parallel Architecture

- **Parallelism at multiple level** is the driving force of computer design
 - Energy and cost being the constraints
- **Classes of parallelism in applications:**
 - Data-Level Parallelism (DLP)
 - Task-Level Parallelism (TLP)
- **Classes of architectural parallelism:**
 - Instruction-Level Parallelism (ILP)
 - Vector architectures/Graphic Processor Units (GPUs)
 - Thread-Level Parallelism
 - Request-Level Parallelism

9

Categories based on Flynn's Taxonomy

- Single instruction stream, single data stream (**SISD**):
 - A sequential computer which exploits instruction level parallelism - uniprocessors
- Single instruction stream, multiple data streams (**SIMD**):
 - A single instruction operates on multiple different data streams
- Multiple instruction streams, multiple data streams (**MIMD**)
 - Multiple autonomous processors simultaneously executing different instructions on different data
- Multiple instruction streams, single data stream (**MISD**)
 - Multiple instructions operate on one data stream
 - No commercial implementation

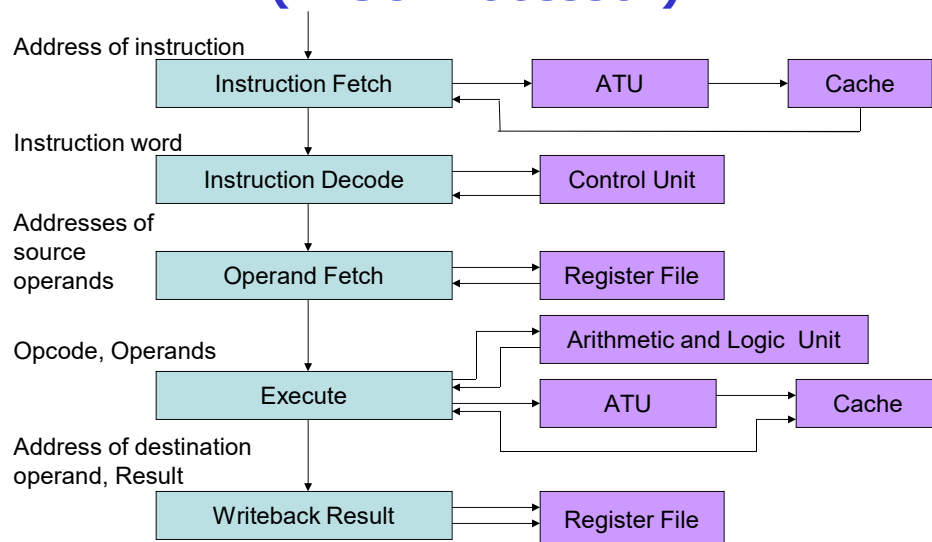
10

Single Instruction Stream, Single Data Stream (SISD)

- Uniprocessor
- Programmer thinks of this as sequential computer
- SISD exploits **instruction-level parallelism (ILP)**
- ILP exploits data level parallelism
 - With compiler help using **pipelining**
 - **Overlapped execution** of instructions, One instruction per clock cycle (**1 IPC**)
 - Using techniques like **superscalar** and **speculative execution**
 - **Instruction-level parallelism**, Multiple instructions per clock cycle (**> 1 IPC**)
- Reason for slowdown in uniprocessor:
 - Diminishing returns in exploiting ILP
 - Growing concern over power

11

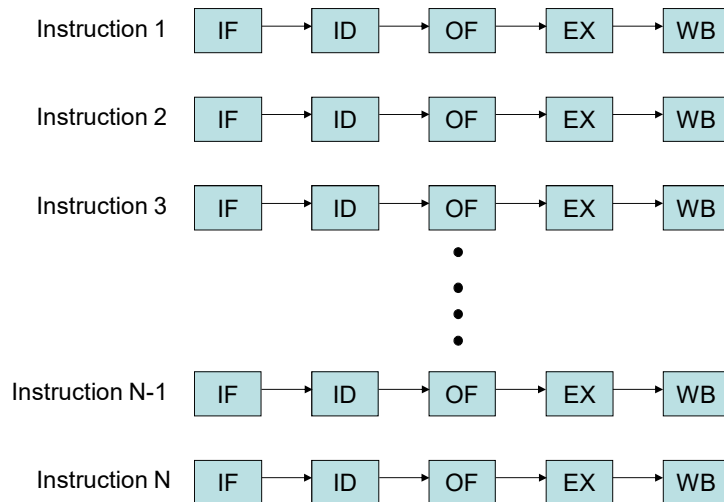
Instruction Cycle in Processors (RISC Processor)



- ATU: Address Translation Unit

12

Non-Pipelined Execution of Instructions




13

Pipelining

- Overlapped execution of instructions
- Different phases of an instruction cycle use different functional units in a processor
- Different phases of multiple instructions can be overlapped
- Target of pipelined processor design
 - Throughput of one instruction per clock cycle

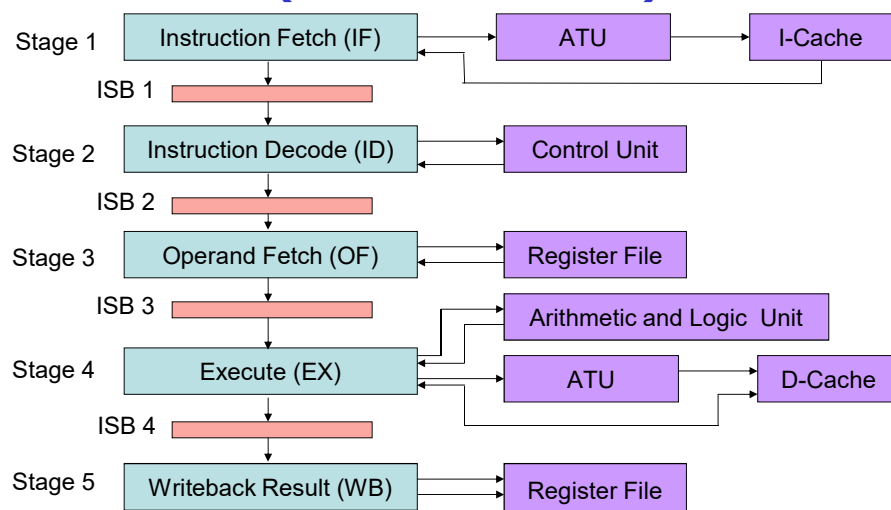
14

Pipelined Execution of Instructions



15

Pipelined Processor Architecture (RISC Processor)



- ISB: Inter Stage Buffer
- ATU: Address Translation Unit
- Split Level I-Cache, Multiport Register File

16


Pipelining Idealism and Realism

- **Pipelining Idealism**
 - **Uniform subcomputations:** Uniform latency
 - **Identical subcomputations:** Fixed latency
 - **Independent subcomputations**
- **Instruction Pipelining Realism**
 - **Non-uniform subcomputations:** Different phases of instruction cycle have different latencies
 - How to decide the period of the pipeline clock cycle?
 - **Balancing the pipeline stages**
 - **Non-identical subcomputations:** A phase of instruction cycle has different latency for different instructions
 - Example: Execute (EX) phase for arithmetic instructions
 - Stalls in the pipeline due to **structural hazards**
 - **Dependent subcomputations:** Source operand of an instruction is the destination operand of a previous instruction
 - Stalls in the pipeline due to **data hazards**

17

Structural Hazards

- Stalls in the pipeline due to **high latency of memory access operations**
 - a cache miss in I-cache during IF phase for I2
 - a cache miss in D-cache during EX phase of I4, a LOAD instruction



- Loss of clock cycles due to these stalls can be reduced by using Level 1 caches with **higher hit rate and lower miss penalty**

18

Structural Hazards

- Stalls in the pipeline due to **arithmetic operations that require multiple clock cycles** during the EX phase of I2 and I4 instructions

I1: ADD R2, R0, R1
 I2: DIV R5, R3, R4
 I3: ADD R11, R9, R0
 I4: FADD F2, F0, F1
 I5: SUB R8, R6, R7

Clock cycle
→

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Instruction	I1	I2	I3	I4	I5									
	IF	ID	OF	EX	WB									
		IF	ID	OF	EX	EX	EX	WB						
			IF	ID	OF			EX	WB					
				IF	ID			OF	EX	EX	EX	EX	WB	
					IF			ID	OF				EX	WB

- Loss of clock cycles due to these stalls can be reduced by using **faster circuits for complex arithmetic operations**

19

Data Hazards

- Stalls in the pipeline due to **Read-After-Write (RAW) data dependencies** among instructions

I1: ADD R2, R0, R1
 I2: SUB R5, R3, R4
 I3: MUL R6, R2, R5; *R2 and R5 are destination operands of I1 and I2*
 I4: ADD R9, R7, R8
 I5: SUB R10, R6, R9; *R6 and R9 are destination operands of I3 and I4*

Clock cycle
→

	1	2	3	4	5	6	7	8	9	10	11	12	13
Instruction	I1	I2	I3	I4	I5								
	IF	ID	OF	EX	WB								
		IF	ID	OF	EX								
			IF	ID	OF								
				IF	ID								
					IF								

- Loss of clock cycles due to these stalls can be reduced by **operand forwarding**

20

Control Hazards

- Stalls in the pipeline due to the presence of **conditional branch instructions**

I1: INC R4
 I2: CMP R4, R5
 I3: BEQ NEXT (Address of I10); *If equal, jump to I10 at NEXT*
 I4:
 I5:
 I6:
 I7:
 I8:
 I9: Instruction
 I10:

	1	2	3	4	5	6	7	8	9	10	11	12
I1:	IF	ID	OF	EX	WB							
I2:		IF	ID	OF	EX	WB						
I3:			IF	ID	OF	EX	WB					
I4:				IF	ID	OF	EX					
I5:					IF	ID	OF					
I6:						IF	ID					
I7:							IF					
I8:												
I9:												
I10:								IF	ID	OF	EX	WB

↑ Flush

- The pipeline is **flushed** in clock cycle 7 because the condition of the branch instruction is TRUE and the branch to the instruction I10 has to take place

21

Static Branch Prediction

- Branch prediction** can be used to reduce the loss of clock cycles due to control hazards
- Branch prediction for conditional branch instructions** due to loop statements can be made by the compiler
 - FOR statement:**
 - Predict that the branch **WILL NOT** take place
 - DO... WHILE statement:**
 - Predict that the branch **WILL** take place
- Instructions are fetched into the pipeline according to the prediction made
 - Whenever the prediction is correct, flushing of the pipeline is avoided
- Prediction is **not possible** for conditional branch instructions due to **IF...THEN...ELSE** statements and **SWITCH** statements.

22

Limitations of Scalar Pipelines

- Scalar pipelines are characterized by a **single-instruction pipeline of K stages**
- All instructions, regardless of type, traverse through the same set of pipeline stages
- At most, one instruction can be resident in each pipeline stage at any one time
- The instructions advance through the pipeline in a **lockstep** fashion
- The maximum throughput for a scalar pipeline is **bounded by one instruction per clock cycle**
- The stalling of a lockstep or rigid pipeline induces unnecessary pipeline bubbles

23


Superscalar Organization

- Target of superscalar processor design
 - **Throughput of multiple instructions per clock cycle**

24

Superscalar Organization

- Target of superscalar processor design
 - Throughput of multiple instructions per clock cycle
- **Parallel pipelines**: Initiate the processing of **multiple instructions** in every clock cycle



- Parallel pipeline of width 4
- 4 instructions can be present in each stage of the pipeline
- **Functional units** in each stage should be **capable of processing 4 instructions** in each clock cycle

- Wider I-cache, Multiport register file, Multiple execute units

25

Superscalar Organization

- **Parallel pipelines**: Initiate the processing of **multiple instructions** in every clock cycle

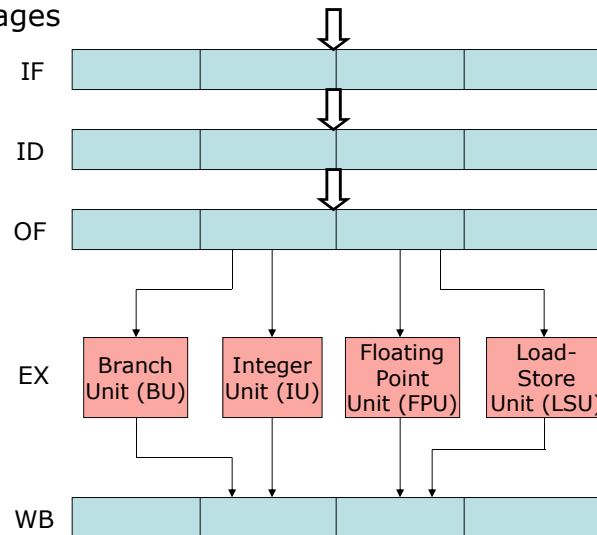
Clock cycle →

		1	2	3	4	5	6	7	8	9	10	11	12	13	14
Instruction ↓	I1	IF	ID	OF	EX	WB									
	I2	IF	ID	OF	EX	WB									
	I3	IF	ID	OF	EX	WB									
	I4	IF	ID	OF	EX	WB									
	I5		IF	ID	OF	EX	WB								
	I6		IF	ID	OF	EX	WB								
	I7		IF	ID	OF	EX	WB								
	I8		IF	ID	OF	EX	WB								
	I9			IF	ID	OF	EX	WB							
	I10			IF	ID	OF	EX	WB							
	I11			IF	ID	OF	EX	WB							
	I12			IF	ID	OF	EX	WB							

26

Superscalar Organization

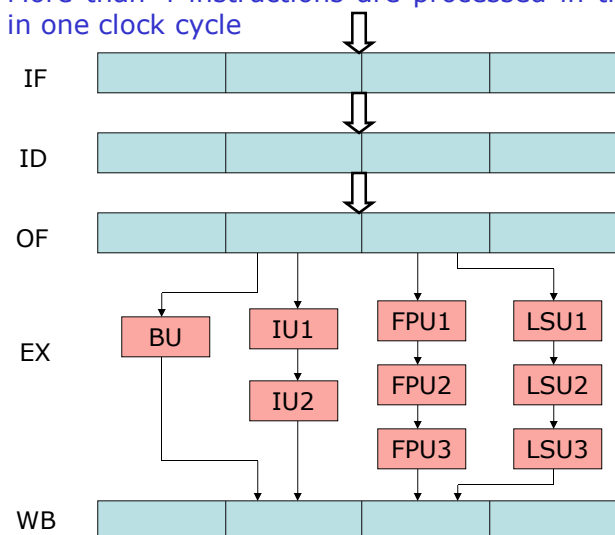
- **Diversified pipelines:** Use multiple and heterogeneous functional units in their Perform Operation (Execute) stages



27

Superscalar Organization

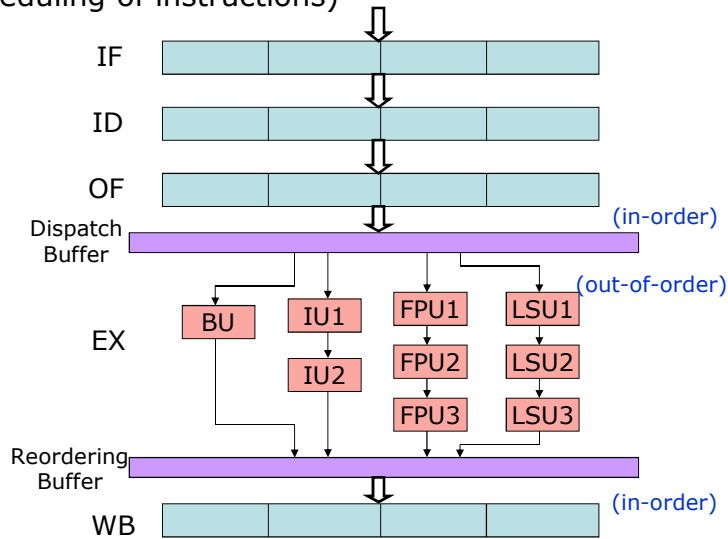
- **Pipelined Execute Units:**
 - More than 4 instructions are processed in the EX stage in one clock cycle



28

Superscalar Organization

- **Dynamic pipelines:** Perform **out-of-order execution** of instructions dynamically during the run-time (Dynamic scheduling of instructions)



29

Superscalar Pipeline Design

- **Instruction Fetch:** **Wide I-cache** is used to fetch multiple instructions in a single read operation
- **Instruction Decode:** RISC instruction set simplifies the decoding task. Multiple instructions can be decoded in a clock cycle
- **Operand Fetch:** Number of read ports for the **multiport register file** is increased
- **Execute:** **One or more execute units of each type** are used. Execute units of some types are also **pipelined**
- **Writeback Result:** Number of write ports for the register file is increased
- Support for **out-of-order execution** and **in-order completion** of instructions is to be provided

30

Data Dependencies

- Consider the following instruction sequence:
 - I1: DIV R8, R0, R1
 - I2: MUL R9, R2, R3
 - I3: SUB R10, R4, R5
 - I4: ADD R8, R8, R9
 - I5: ADD R6, R8, R10
 - I6: MUL R9, R0, R2
 - I7: SUB R7, R9, R3
- True data dependencies (RAW dependencies):
 - I4 \rightarrow I1: R8
 - I4 \rightarrow I2: R9
 - I5 \rightarrow I3: R10
 - I5 \rightarrow I4: R8
 - I7 \rightarrow I6: R9
- False data dependencies (WAW and WAR dependencies):
 - I4 \rightarrow I1: R8 WAW (Output dependence)
 - I6 \rightarrow I2: R9 WAW (Output dependence)
 - I6 \rightarrow I4: R9 WAR (Anti-dependence)

31

Dependency Detection and Resolution

- A 'valid'(V) bit is associated with each register
- RAW dependency
 - Let R_i be the destination operand of an instruction I_m
 - The 'valid' bit of R_i is set to 0 while decoding I_m
 - The 'valid' bit of R_i is set to 1 when the Writeback for I_m is done, i.e., when I_m is completed
 - Let R_j be the source operand of an instruction I_n with $n > m$
 - If the 'valid' bit of R_i is 1, take the contents of R_i
 - If the 'valid' bit of R_i is 0, then it is a case of RAW dependency
 - The instruction I_n will wait in the reservation station till the result of I_m is available on the operand forwarding path

32

Dependency Detection and Resolution

- **WAW dependency (Output dependency)**
 - Let R_i be the **destination operand** of an instruction I_n with $n > m$
 - If the 'valid' bit of R_i is 1, there is no WAW dependency
 - If the 'valid' bit of R_i is 0, then it is a case of **WAW dependency**
 - The register R_i is renamed
- **WAR dependency (Anti-dependency)**
 - There is no need to check for it when the instructions are dispatched and completed in-order

33

Register Renaming Techniques

- Register renaming using Rename Registers
 - A non-architectural register file called **Rename Register File (RRF)** is used
- **Method 1:** Rename the destination operand of every instruction
 - It requires large number of rename registers
- **Method 2:** Rename the destination operand of an instruction that has WAW dependency
 - It is necessary to check for WAW dependency while decoding an instruction

34

Register Renaming

- Registers in RRF are RR_0, RR_1, \dots, RR_M
- Consider the following instruction sequence:


```

I1: DIV    R8, R0, R1
I2: MUL    R9, R2, R3
I3: SUB    R10, R4, R5
I4: ADD    R8, R8, R9
I5: ADD    R6, R8, R10
I6: MUL    R9, R0, R2
I7: SUB    R7, R9, R3
      
```
- False data dependencies (WAR and WAW dependencies):


```

I4  → I1: R8  WAW
I6  → I2: R9  WAW
      
```
- Instruction sequence of register renaming


```

I4: ADD    RR0, R8, R9
I5: ADD    R6, RR0, R10
I6: MUL    RR1, R0, R2
I7: SUB    R7, RR1, R3
      
```

35

Multiple Instruction Stream, Multiple Data Stream (MIMD)

- Multiprocessor:
 - Computers consists of tightly coupled processors
 - Coordination and usage are controlled by a single operating system
 - Share memory through a shared address space
- Each processor fetches its own instructions and operate on its own data
- MIMD targets task-level parallelism as well as data-level parallelism
- Types of MIMD architectures:
 - Tightly-coupled MIMD
 - All the processing elements have shared memory model.
 - Loosely-coupled MIMD
 - All the processing elements have separate memory model

36

Tightly Coupled MIMD Architecture

- Exploits **thread-level parallelism**
- Thread is a light-weight process
- Threads of a process share the **same virtual address space**
- Multiple cooperating threads operate in parallel
- **Parallel processing**: Tightly coupled set of threads collaborating on a single task
- This architecture ranges from dual processor to dozens of processors
 - Communicate and coordinate through the sharing memory
- Multiprocessors include in
 - Computers consists of **single chip** with multiple cores (multicore)
 - Computers consisting of **several chips**, each may be multicore design

37

Loosely Coupled MIMD Architecture

- Exploits **request-level parallelism**
 - Exploits parallelism among largely **decoupled tasks** specified by programmer or the operating system
- Many **independent tasks processed in parallel** with little need for communication and synchronization
- Examples: **Clusters** and Warehouse-scale computers (**Cloud computers**)

38

Single Instruction Stream, Multiple Data Stream (SIMD)

- Same instruction is executed by multiple processors using different data streams
- SIMD exploits **data-level parallelism** by applying the same operation to multiple items of data in parallel
- Data parallelism for
 - Matrix-oriented scientific computing
 - Media-oriented image and sound processors
- Each processor has its own data memory, but there is a single instruction memory and control processor
- **SIMD vs MIMD**:
 - MIMD is more flexible and expensive than SIMD
 - SIMD is more energy efficient than MIMD
 - In SIMD programmer continues to think sequentially, yet achieves parallel speedup

39

Single Instruction Stream, Multiple Data Stream (SIMD)

- **Types of SIMD architecture**:
 - Vector architectures
 - Multimedia extensions to standard instruction sets
 - Graphics processor units (GPUs)
- For x86 processors:
 - Expect two additional cores per chip per year
 - Potential speedup from SIMD to be twice that from MIMD!

40

Detecting and Enhancing Loop-Level Parallelism

- This is critical for exploiting Data-Level Parallelism and Task-Level Parallelism
- Compiler technology for discovering the amount of parallelism that we can exploit in a program as well as hardware support for these compiler techniques
- **Focus** is on understanding
 - When a loop is parallel
 - How dependence can prevent a loop from being parallel
 - Techniques for eliminating some types of dependences
- Loop-level parallelism is normally analyzed at the source level or close to it
- Loop-level analysis involves **determining what dependences exist among the operands** in a loop across the iterations of the loop
- We focus on **data dependences** (RAW dependence)

41

Loop-Carried Dependence

- Dependence exists among the operand (data value) in one iteration with the data value produced in earlier iterations
- **Loop-carried dependence generally prevent parallelism**
- **Example:**

```
for (i=999; i>=0; i--)
    x[i] = x[i] + s;
```
- Loop-carried dependence between **successive uses of induction variable i in different iterations**
 - Handled using **loop unrolling** techniques
- Finding loop-level parallelism involves recognizing structures such as
 - Loops
 - Array references
 - Induction variable compositions

42

Loop-Carried Dependence

- **Example:**

```
for (i=0; i< 100; i++){
    A[i+1] = A[i] + C[i];    // S1
    B[i+1] = B[i] + A[i+1]; // S2
}
```

- Statement S1 and S2 has loop-carried dependence
- Statement S2 is also depending on S1 within an iteration
 - This is **intra-loop dependence**

43

Loop-Carried Dependence

- **Example:**

```
for (i=0; i< 100; i++){
    A[i] = A[i] + B[i];    // S1
    B[i+1] = C[i] + D[i]; // S2
}
```

- Loop-carried dependence between S2 and S1
- Loop is parallel if it can be written without a cycle in dependence
- Absence of a cycle means that the **dependence give a partial ordering** on the statements
- It must be **transformed** to confirm to the partial ordering and expose parallelism
- Two observations:
 - There is no dependence from S1 and S2
 - On the first iteration of loop, S2 depends on the value B[0] computed prior to the initiating the loop

44

Loop-Carried Dependence

- Replace with code sequence

```
A[0]=A[0]+B[0];
for (i=0; i< 99; i++){
    B[i+1] = C[i] + D[i];
    A[i+1] = A[i+1] + B[i+1];
}
B[100]=C[99] + D[99];
```

- **Example:**

```
for (i=0; i< 100; i++){
    A[i] = B[i] + C[i];    // S1
    D[i] = A[i] * E[i];    // S2
}
```

45

Loop-Carried Dependence

- Often loop-carried dependence are in the form of **recurrence**
- Detecting recurrence (loop-carried dependence) is important as some architectures have special support for executing recurrences

46

Name Dependence

- Name dependence:
 - Eliminated by renaming and copying
- Example:

```
for (i=0; i< 100; i++){
    y[i] = x[i] / c;    // S1
    x[i] = x[i] + c;    // S2
    z[i] = y[i] + c;    // S3
    y[i] = c - y[i];    // S4
}
```

- True dependence
- Anti-dependence
- Output dependence

47

Eliminating Dependent Computations

- Recurrences
- Example: Dot product


```
for (i=0; i< 1000; i++){
    sum = sum + x[i]*y[i];
}
```
- Loop-carried dependence on variable `sum`
- Transform into a set of loops, one is completely parallel and another is partial
 - **Scalar expansion**: This transformation makes loop completely parallel
 - **Reduce step**: Sums up the elements of vector
 - Reductions are handled by special hardware in vector and SIMD architectures
 - It allows reduction step to be done much faster than it could be done in scalar mode

48

Dependence Analysis

- Dependence Analysis is critical for exploiting parallelism
- For detecting loop-level parallelism dependence analysis is a basic tool
- However, it applies only under a limited circumstances
- CUDA/OpenMP programmers write explicitly parallel loop

49

Text Books

- J. L. Hennessy and D. A. Patterson, *Computer Architecture – A Quantitative Approach*, 5th Edition, Morgan Kaufmann, 2012
- J. P. Shen and M. H. Lipasti, *Modern Processor Design – Fundamentals of Superscalar Processors*, Tata McGraw-Hill, 2005.
- W. Stallings, *Computer Organization and Architecture – Designing for Performance*, 8th Edition, Prentice Hall of India, 2011.
- D. Sima, T. Fountain and P. Kacsuk, *Advanced Computer Architecture – A Design Space Approach*, Addison-Wesley, 1997.

50