

Parallel Architecture **[High Performance (HPC) Architecture]**

Categories based on Flynn's Taxonomy

- Single instruction stream, single data stream (SISD):
 - A sequential computer which exploits no parallelism in either the instruction or data streams
 - One instruction pool acting on one data pool

Categories based on Flynn's Taxonomy

- Single instruction stream, multiple data streams (SIMD):
 - A single instruction operates on multiple different data streams
 - SIMD exploits data-level parallelism by applying the same operation to multiple items of data in parallel
 - Data parallelism for
 - Matrix-oriented scientific computing
 - Media-oriented image and sound processors
 - Each processor has its own data memory, but there is a single instruction memory and control processor
 - Types of SIMD architecture:
 - Vector (array) processors
 - Graphics processor units (GPUs)

3

Categories based on Flynn's Taxonomy

- Multiple instruction streams, multiple data streams (MIMD)
 - Multiple autonomous processors simultaneously executing different instructions on different data
 - Types of MIMD architectures:
 - Tightly-coupled MIMD: All the processing elements have shared memory model
 - Loosely-coupled MIMD: All the processing elements have separate memory model
 - Multiprocessors
 - Computers consists of single chip with multiple cores (multicore)
 - Computers consisting of several chips, each may be multicore design

4



Multiprocessors: Parallel Architecture

- Shared memory architecture
- Distributed memory architecture
- Hybrid architecture

5

Shared Memory Architecture


- **Centralised multiprocessors:**
 - Multiple processors attached to a bus
 - all the processors share the same primary memory
 - **Uniform Memory Access (UMA)** or **Symmetric Multiprocessor (SMP)**
- **Distributed multiprocessors:**
 - Processors are distributed and connected by bus interconnect
 - Distributed collection of memories forms one logical address space
 - Same address on different processors refers to the same memory location
 - **Nonuniform Memory Access (NUMA)**

Support of **Cache Coherence**

6

Distributed Memory Architecture


- **Distributed multiprocessors:**
 - Each processor has its local memory with the address space available only for this processor
 - Processors can exchange data through the **interconnection network** by means of communication by **passing messages**
 - Required a communication network to connect inter-processor memory
 - Each processor operated independently
 - No concept of cache coherency



7

Hybrid Memory Architecture

- **Distributed multiprocessors:**
 - Employ both shared and distributed memory architectures
 - Shared memory could be **shared memory machine** or with **GPUs**



8


Parallel Programming Models

- **Shared memory:** All tasks have visibility to a global address space
 - **OpenMP**
- **Distributed memory:** All tasks have their own private address space and communicate via messages
 - **Message Passing Interface (MPI)**


9

Parallel Programming Models

- **Hybrid memory:** Combination of shared and distributed memory
 - **OpenMP + MPI**



- **CUDA + MPI**




10

Parallel Algorithm Design

Task/Channel Model


- Described by Ian Foster [1]
- Facilitates the development of **efficient parallel algorithms**, particularly distributed memory architecture
- It represents a parallel computation as a **set of tasks** that may interact with each other by **sending messages through channels**



[1] Ian Foster, *Designing and Building Parallel Programs: Concepts and Tools for Parallel Software Engineering*, Addison-Wesley, 1995.

Task/Channel Model


- **Task:** A task is a program, its local memory and a collection of I/O ports
 - The local memory contains the program's instruction and its private data
 - A task can send local data values to other tasks via output ports
 - A task can receive local data values from other tasks via input ports



13

Task/Channel Model


- **Task:** A task is a program, its local memory and a collection of I/O ports
 - The local memory contains the program's instruction and its private data
 - A task can send local data values to other tasks via output ports
 - A task can receive local data values from other tasks via input ports



14

Task/Channel Model

- **Task:** A task is a program, its local memory and a collection of I/O ports



- **Channel:** Message queue that connects one task's output port with another task's input port

15

Foster's Design Methodology


- Four-step process for designing parallel algorithm [1]
 - Partitioning
 - Communication
 - Agglomeration
 - Mapping

[1] Ian Foster, *Designing and Building Parallel Programs: Concepts and Tools for Parallel Software Engineering*, Addison-Wesley, 1995.

16

Partitioning

- Process of dividing the computation and data into pieces

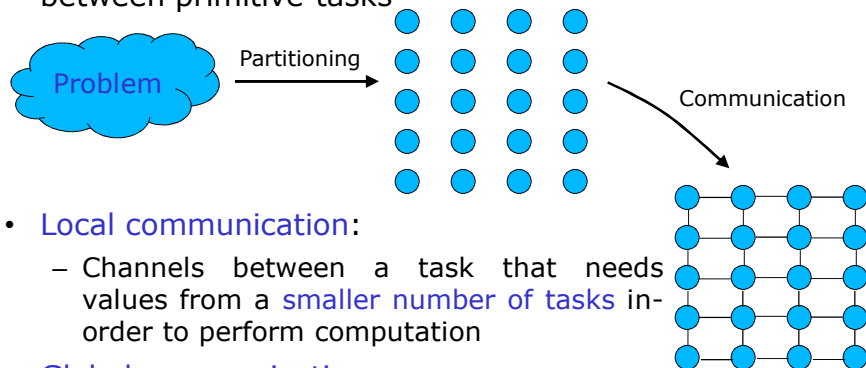


- **Domain decomposition (data-centric):**
 - First divide the data into pieces and then determine how to associate computations with the data
 - Focus on **largest or frequently accessed data structure**
- **Functional decomposition (computation-centric):**
 - First divide the computation into pieces and then determine how to associate data item into individual computations
 - Collection of tasks and running concurrently
- **Primitive task**

17

Communication

- Process of **determining the communication pattern** between primitive tasks

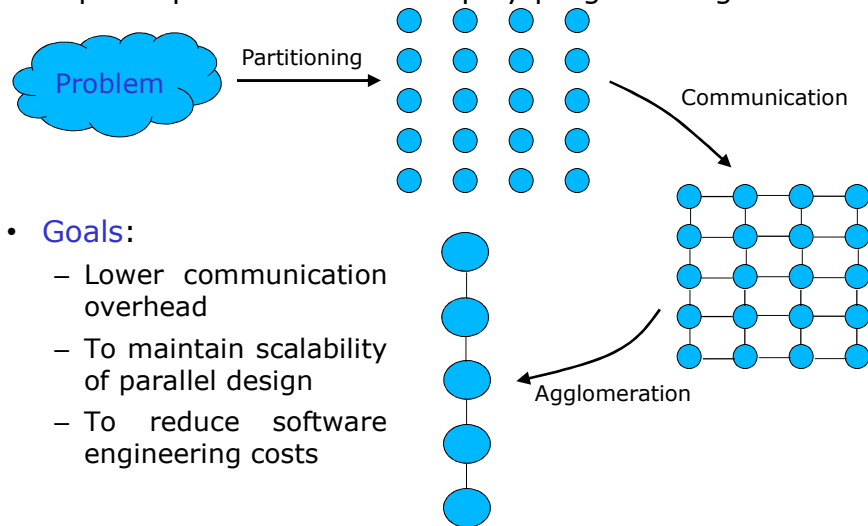


- **Local communication:**
 - Channels between a task that needs values from a **smaller number of tasks** in-order to perform computation
- **Global communication:**
 - Channels between a task and a **significant number of primitive tasks** in-order to perform computation

18

Agglomeration

- Process of **grouping tasks into larger tasks** in-order to improve performance or simplify programming

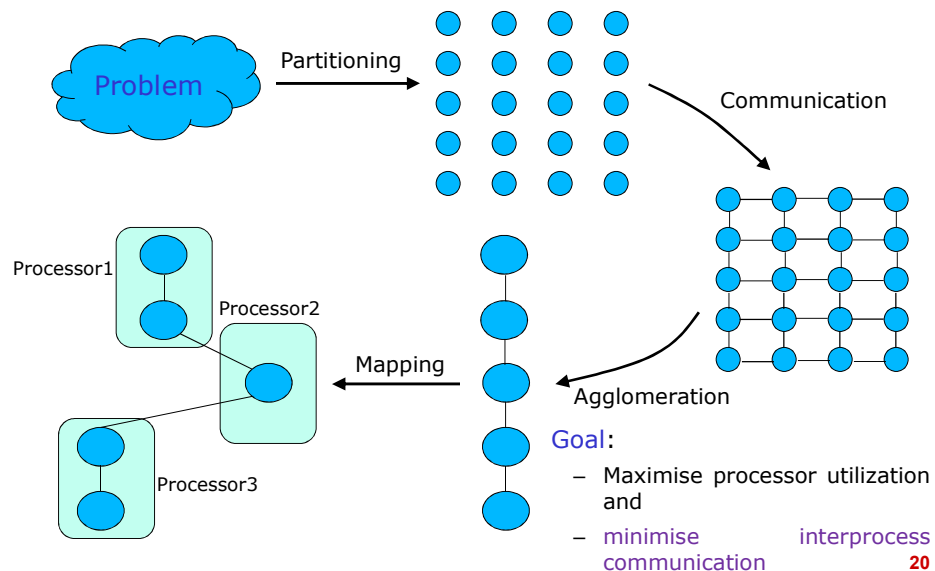


- Goals:**
 - Lower communication overhead
 - To maintain scalability of parallel design
 - To reduce software engineering costs

19

Mapping

- Process of **assigning tasks to processors**



Goal:

- Maximise processor utilization and
- minimise interprocess communication

20

Example: Finding Minimum in a Matrix

21

Text Books

- M. J. Quinn, *Parallel Programming in C with MPI and OpenMP*, McGraw-Hill, 2004.
- Ian Foster, *Designing and Building Parallel Programs: Concepts and Tools for Parallel Software Engineering*, Addison-Wesley, 1995.

22