


# **Message Passing Interface (MPI)**

## **Parallel Programming Approaches**

- **Approach1**: Parallelising compilers and high-level parallel programming languages
- **Approach2**: Augmenting an existing sequential language with low level constructs expressed by functional calls or compiler directives
  - Parallel programming in C, C++ and Fortran with MPI and OpenMP

## Message-Passing Model

- Collection of processors, each with its own memory
  - Processor has direct access only to the instruction and data stored in local memory
- Interconnection network supports message passing between the processors
- User specifies the number of concurrent processes when the program begins
- Every process executes the same program
- Each process has unique ID number



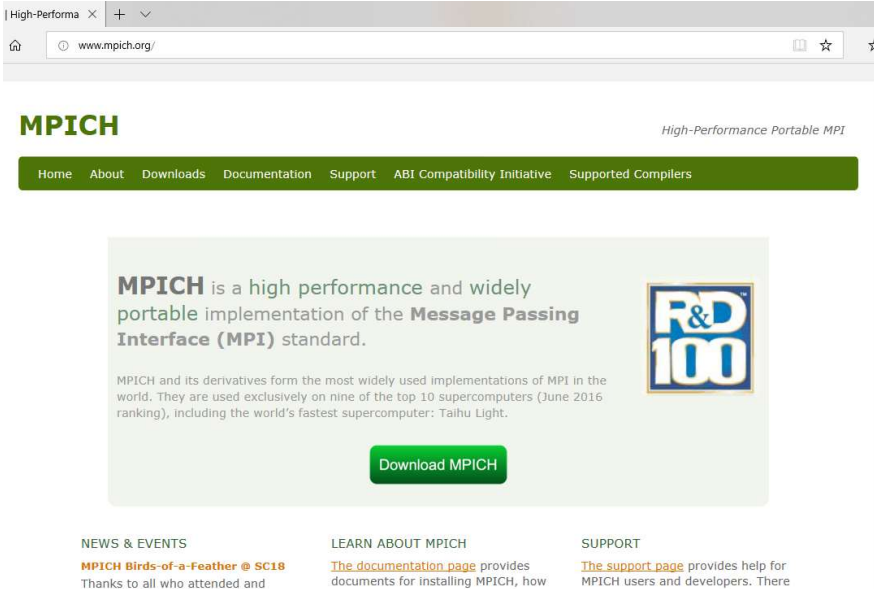
3

## MPI Standards and Goals

- MPI-1 standard defined in 1994
- Current standard is MPI-3
- Standard gives the **uniformity to the code written** so that code can be **portable**, **compile** and **run on any platform that support MPI standard**
- Standard specifies the **names**, **calling sequence** and **subroutines and functions** called from C, C++ and Fortran
- **Goals:**
  - To Provide source code portability
  - To allow efficient implementations across a range of architectures
  - To support heterogeneous parallel architecture

4

## Installing MPICH3 in a Single Machine

A screenshot of the MPICH website. The browser address bar shows 'www.mpich.org/'. The page has a green navigation bar with links: Home, About, Downloads, Documentation, Support, ABI Compatibility Initiative, and Supported Compilers. The main content area features the MPICH logo, a description of it as a high-performance and widely portable implementation of the Message Passing Interface (MPI) standard, and a 'Download MPICH' button. To the right is an 'F&D 100' logo. Below the main content are three sections: 'NEWS & EVENTS' with a link to 'MPICH Birds-of-a-Feather @ SC18', 'LEARN ABOUT MPICH' with a link to 'The documentation page', and 'SUPPORT' with a link to 'The support page'.

5

## Installing MPICH3 in a Single Machine

```
>> tar -xzf mpich-3.3.tar.gz

>> cd mpich-3.3

>> ./configure --disable-fortran

>> sudo make install

>> mpiexec --version
```

---

<https://www.mpi-forum.org/docs/>

6

## MPI Programming

- The program begins with pre-processor directive to include header file for MPI

```
#include <stdio.h>
```

```
#include <mpi.h>
```

```
void main (int argc, char *argv[]){
```

```
}
```

- `argc` and `argv` are needed to pass to the function that initialises MPI

- Some of the **MPI function calls**:

- `MPI_Init`: To initialize MPI
- `MPI_Comm_rank`: To determine process's ID number
- `MPI_Comm_size`: To find the number of processes
- `MPI_Finalize`: To shut down the MPI

7

## MPI Overview

- MPI initialization and finalization
- MPI communication environment, size and rank
- MPI datatypes
- Point-to-point communication
  - MPI send and receive
  - Blocking and deadlock
- Collective communication
  - Barrier
  - Broadcast
  - Reduction
  - Gather
  - Scatter
- Time functions

8

## Function `MPI_Init`

- The first MPI function call made by every MPI process
- It allows the system to do any startup needed to handle further calls to the MPI library

```
MPI_Init (&argc, &argv);
```

## Function `MPI_Finalize`

- The last MPI function call made by every MPI process i.e. called after a process has completed all its MPI library calls
- It allows the system to free up resources such as memory that has been allocated to MPI

```
MPI_Finalize ();
```

9

## MPI Program Compile and Execution

- Compiling:
  - `mpicc -o example.out example.c`
- Execution:
  - `mpiexec -n <num of processes> ./example.out`

10

## MPI Datatypes

- MPI constants associated with C datatypes:

Name	C datatype
MPI_CHAR	signed char
MPI_DOUBLE	double
MPI_FLOAT	Float
MPI_INT	int
MPI_LONG	long
MPI_LONG_DOUBLE	long double
MPI_SHORT	short
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long
MPI_UNSIGNED_SHORT	unsigned short

11

## Special MPI Datatypes for C

- Special MPI datatypes include:

Name	Meaning
MPI_Comm	Communicator
MPI_Status	Structure containing several pieces of status information for MPI calls

- **Communicator**: An opaque object that provide the environment for message passing among processes
- When MPI is initialised, every active process becomes a member of a communicator
  - MPI\_COMM\_WORLD: an MPI constant – a default communicator

12

## Function MPI\_Comm\_rank

- Processes within a communicator are ordered
- **Rank**: Position of a process in the overall order
- Each process has a unique rank (ID number)
- Function to **determine rank of processes within a communicator**

```
int MPI_Comm_rank (MPI_Comm comm, int *rank);
```

## Function MPI\_Comm\_size

- To determine total number of processes in a communicator

```
int MPI_Comm_size (MPI_Comm comm, int *size);
```

13

## MPI Overview

- MPI initialization and finalization
- MPI communication environment, size and rank
- MPI datatypes
- Point-to-point communication
  - MPI send and receive
  - Blocking and deadlock
- Collective communication
  - Barrier
  - Broadcast
  - Reduction
  - Gather
  - Scatter
- Time functions

14

## Point-to-Point Communication

- One process send message and other process receive that message



- For point-to-point communication we need to specify **source** and **destination**
- Any message that is passed will have **envelope** and **message body**

15

## Point-to-Point Communication

- **Envelope**: Contain source and destination address
  - Source : The sending process
  - Destination : The receiving process
  - Communicator: Specifies a group of processes to which both sender and receiver belongs
  - tag : used to classify the messages
- **Message body**:
  - Buffer : The message data (Buffer as an array)
  - Datatype : The type of message data
  - Count : The number of items of type datatype in buffer

16



## Function MPI\_Send

```
int MPI_Send (void *buff, int count,
MPI_Datatype dtype, int dest, int tag, MPI_Comm
comm);
```

- Blue coloured arguments are message body
- Red coloured arguments are envelope
- All arguments are input arguments
- An error code is returned by the function

## Function MPI\_Recv

```
int MPI_Recv (void *buff, int count,
MPI_Datatype dtype, int source, int tag,
MPI_Comm comm, MPI_Status *status);
```

- buff and status are output arguments
- Rest are all input arguments

17

## Completion and Blocking in Point-to-Point Communication

- Completion for a call MPI\_Recv:
  - Matching message has arrived and the message data have been copied into the output argument of the call
- Completion for a call MPI\_Send:
  - Message specified in the input argument of the call has been handed off to MPI
    - Variable passed as input argument can now be over written and reused
    - The message is copied to internal buffer for later delivery
- MPI will wait for the destination process to receive the message
- Blocking:
  - When the message passed through MPI\_Send is larger than the available buffer at receiving side, sending process will be blocked till the receiving process start receiving or sufficient buffer is available

18

## Deadlock and Resolving Deadlock

- **Deadlock:**
  - Occurs when two or more processes are blocked
  - Each is waiting for the other to make progress
- **Example:**
  - Process 0 cannot proceed until process 1 sends a message
  - Process 1 cannot proceed until process 0 sends a message
- **Resolving Deadlock:**
  - Process 0 receives from process 1 and then sends the message to process 1
  - Process 1 sends message to process 0 and then received from process 0
- **Sending the message depends on the available buffer in the receiving side**

19


## MPI Overview

- MPI initialization and finalization
- MPI communication environment, size and rank
- MPI datatypes
- Point-to-point communication
  - MPI send and receive
  - Blocking and deadlock
- Collective communication
  - Barrier
  - Broadcast
  - Reduction
  - Gather
  - Scatter
- Time functions

20

## Collective Communication

- Communication operation in which a group of processes work together to distribute or gather together a set of one or more values
- One-to-many or many-to-one communication
- **Barrier Synchronization:**
  - Synchronising all the processes in the communicator




- **Function MPI\_Barrier:**

```
int MPI_Barrier (MPI_Comm comm);
```

21

## Broadcast

- Enables a process to broadcast one or more data items of the same type to all other processes in the communicator




- **Function MPI\_Bcast:**

```
int MPI_Bcast (
    void          *buffer, // Address of first data item to be broadcasted
    int           count,    // number of elements to be broadcasted
    MPI_Datatype  dtype,    // Type of element to be broadcasted
    int           root,     // ID (rank) of the process doing broadcast
    MPI_Comm      comm);    // Communicator
```

22

## Reduction

- Performs **one or more reduction operations** on the values submitted from all the processes in a communicator



- Function MPI\_Reduce:**

```
int MPI_Reduce (
    void          *send_buffer, // Address of first reduction element
    void          *recv_buffer, // Address of reduction result
    int           count,        // number of elements to be send or reduction
    MPI_Datatype  dtype,        // Type of element
    MPI_Op        Operator,     // Reduction operator
    int           root,         // ID (rank) of the process receiving
    MPI_Comm      comm);       // Communicator
```

23

## Reduction Operators


- MPI's built-in reduction operators:

Name	Meaning
MPI_BAND	Bitwise and
MPI_BOR	Bitwise or
MPI_BXOR	Bitwise exclusive or
MPI_LAND	Logical and
MPI_LOR	Logical or
MPI_LXOR	Logical exclusive or
MPI_MAX	Maximum
MPI_MAXLOC	Maximum and its location
MPI_MIN	Minimum
MPI_MINLOC	Minimum and its location
MPI_PROD	Product
MPI_SUM	Sum

24

## Gather

- Enables each process in a communicator to **send the contents of its send buffer to the root process**
- Root process receives the messages and store them in rank order




### Function MPI\_Gather:

```
int MPI_Gather (
    void          *send_buffer, // Starting address of send buffer (IN)
    int           send_count,   // number of elements in send buffer (IN)
    MPI_Datatype  send_dType,   // Type of send buffer element (IN)
    void          *recv_buffer, // Starting address of receiving buffer (OUT)
    int           recv_count,    // number of elements in recv buffer (IN)
    MPI_Datatype  recv_dType,    // Type of receiving buffer element (IN)
    int           root,         // ID (rank) of the process doing gathering (IN)
    MPI_Comm      comm);       // Communicator (IN)
```

25

## Scatter

- A group of elements held by the root process is **divided into equal chunks** and **one chunk is send to every process** in the communicator including root.



### Function MPI\_Scatter:

```
int MPI_Scatter (
    void          *send_buffer, // Starting address of send buffer (IN)
    int           send_count,   // number of elements in send by each process (IN)
    MPI_Datatype  send_dType,   // Type of send buffer element (IN)
    void          *recv_buffer, // Starting address of receiving buffer (OUT)
    int           recv_count,    // number of elements in recv buffer (IN)
    MPI_Datatype  recv_dType,    // Type of receiving buffer element (IN)
    int           root,         // ID (rank) of the process doing gathering (IN)
    MPI_Comm      comm);       // Communicator (IN)
```

26

## MPI Overview

- MPI initialization and finalization
- MPI communication environment, size and rank
- MPI datatypes
- Point-to-point communication
  - MPI send and receive
  - Blocking and deadlock
- Collective communication
  - Barrier
  - Broadcast
  - Reduction
  - Gather
  - Scatter
- **Time functions**

27

## Time Functions for Measuring Performance

- One way to measure the performance of a parallel application is to look at the wall clock time.
- It measures the number of seconds that elapse from the time we initiate the execution until the program terminates
- **Function MPI\_Wtime:**
  - Returns the number of seconds that have elapsed since some point of time in the past

```
double MPI_Wtime (void);
```
- **Function MPI\_Wtick:**
  - Returns the precision of the result returned by MPI\_Wtime

```
double MPI_Wtick (void);
```

28

## Example: Innerproduct Between Two Vector

- Let  $\mathbf{x}=[x_1, x_2, \dots, x_d]^\top$  and  $\mathbf{y}=[y_1, y_2, \dots, y_d]^\top$  be two vectors
- Innerproduct between two vector is given as:

$$z = \mathbf{x}^\top \mathbf{y}$$
$$z = \sum_{i=1}^d (x_i y_i)$$

- Assignment:
  - Read or Create the vectors in root processor (i.e. id=0)
  - Use scatter function
  - Use reduce function

29

## Text Books

- <https://www.mpi-forum.org/docs/>
- M. J. Quinn, *Parallel Programming in C with MPI and OpenMP*, McGraw-Hill, 2004.

30